

ShritejShrikant_file2_hw2

September 10, 2023

0.1 Install/Import *Libraries*

[3]:

100.3/100.3

kB 2.1 MB/s eta 0:00:00

[43]:

```
# Import necessary libraries
import pandas as pd
from pathlib import Path

# Import the joblib library for saving and loading models
import joblib

# Import scikit-learn classes for building models
from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import *
from sklearn.metrics import classification_report
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.base import TransformerMixin, BaseEstimator

from skopt.space import Real, Categorical, Integer
from sklearn.metrics import precision_recall_curve, auc, make_scorer,
    ↪cohen_kappa_score, balanced_accuracy_score, confusion_matrix
from skopt import BayesSearchCV
from skopt.space import Real, Categorical, Integer

import spacy

# Import the scipy library for working with sparse matrices
from scipy.sparse import csr_matrix
```

[5]:

```
import sys
if 'google.colab' in str(get_ipython()):
```

```

from google.colab import drive
drive.mount('/content/drive')

!pip install -U nltk -qq
!pip install -U spacy -qq
!python -m spacy download en_core_web_sm -qq
!pip install -U scikit-optimize -qq

basepath = '/content/drive/MyDrive/NLP/'
sys.path.append('/content/drive/MyDrive/NLP/custom-functions')
else:
basepath = '/home/harpreet/Insync/google_drive_shaannoor/data'
sys.path.append(
    '/home/harpreet/Insync/google_drive_shaannoor/data/custom-functions')

```

Mounted at /content/drive
2023-09-10 17:50:56.766502: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not find TensorRT

12.8/12.8 MB

103.9 MB/s eta 0:00:00

Download and installation successful

You can now load the package via `spacy.load('en_core_web_sm')`

[6]: `sys.path`

```

[6]: ['/content',
      '/env/python',
      '/usr/lib/python310.zip',
      '/usr/lib/python3.10',
      '/usr/lib/python3.10/lib-dynload',
      '',
      '/usr/local/lib/python3.10/dist-packages',
      '/usr/lib/python3/dist-packages',
      '/usr/local/lib/python3.10/dist-packages/IPython/extensions',
      '/root/.ipython',
      '/content/drive/MyDrive/NLP/custom-functions']

```

```

[7]: base_folder = Path(basepath)
      data_folder = base_folder/'datasets/spam'
      model_folder = base_folder/'models/spam'
      custom_functions = base_folder/'custom-functions'

```

```

[8]: import custom_preprocessor_mod as cp
      from featurizer import ManualFeatures
      from plot_learning_curve import plot_learning_curve

```

0.2 Load Dataset

Downloaded the dataset from here : <https://www.kaggle.com/datasets/uciml/sms-spam-collection-dataset>

```
[9]: data = pd.read_csv(data_folder/'spam.csv', encoding='latin-1')
data.head()
```

```
[9]:      v1                                     v2 Unnamed: 2  \
0  ham  Go until jurong point, crazy.. Available only ...      NaN
1  ham                                     Ok lar... Joking wif u oni...      NaN
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...      NaN
3  ham  U dun say so early hor... U c already then say...      NaN
4  ham  Nah I don't think he goes to usf, he lives aro...      NaN

      Unnamed: 3 Unnamed: 4
0           NaN          NaN
1           NaN          NaN
2           NaN          NaN
3           NaN          NaN
4           NaN          NaN
```

```
[10]: data.shape
```

```
[10]: (5572, 5)
```

```
[11]: data.isnull().sum()
```

```
[11]: v1           0
v2           0
Unnamed: 2    5522
Unnamed: 3    5560
Unnamed: 4    5566
dtype: int64
```

```
[12]: data['v1'].value_counts()
```

```
[12]: ham      4825
spam      747
Name: v1, dtype: int64
```

```
[13]: data.drop(columns=['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], inplace=True)
data.rename(columns={'v1': 'label', 'v2': 'text'}, inplace=True)
data.head()
```

```
[13]:   label                                     text
0  ham  Go until jurong point, crazy.. Available only ...
1  ham                                     Ok lar... Joking wif u oni...
```

```

2 spam Free entry in 2 a wkly comp to win FA Cup fina...
3 ham U dun say so early hor... U c already then say...
4 ham Nah I don't think he goes to usf, he lives aro...

```

```
[14]: data['label'].value_counts(normalize=True)*100
```

```

[14]: ham      86.593683
      spam     13.406317
      Name: label, dtype: float64

```

```
[14]:
```

```
[15]: # prompt: convert above labels: 0 for 'ham' & 1 'spam for above dataset
```

```

data['label'].replace(['ham', 'spam'], [0, 1], inplace=True)
data['label'].value_counts(normalize=True)*100

```

```

[15]: 0      86.593683
      1      13.406317
      Name: label, dtype: float64

```

```

[16]: print(data.shape)

print(data['label'].value_counts())

x = data['text']
y = data['label']

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
↳ random_state=21, stratify=y, shuffle=True)

print(x_train.shape, x_test.shape, y_train.shape, y_test.shape)

print(y_train.value_counts())
print(y_test.value_counts())

```

```

(5572, 2)
0      4825
1       747
Name: label, dtype: int64
(4457,) (1115,) (4457,) (1115,)
0      3859
1       598
Name: label, dtype: int64
0       966
1       149

```

Name: label, dtype: int64

```
[17]: xtrain = x_train.values
      ytrain = y_train.values

      xtest = x_test.values
      ytest = y_test.values

      print(xtrain.shape, ytrain.shape)
```

(4457,) (4457,)

```
[17]:
```

0.3 Spacy

```
[18]: # Spacy Tokenizer
      # Loading the 'en_core_web_sm' language model from the spaCy library
      nlp = spacy.load('en_core_web_sm')

      disabled = nlp.select_pipes(
          disable=['tok2vec', 'tagger', 'parser', 'attribute_ruler', 'lemmatizer',
                  ↪'ner'])

      def spacy_tokenizer(data):
          doc = nlp(data)
          return [token.text for token in doc]

      cpp = cp.SpacyPreprocessor(
          model='en_core_web_sm')

      def spacy_preprocessor(text):
          filtered_text = cpp.transform([text])
          return " ".join(filtered_text)
```

```
[19]: # save this to a file
      X_train_cleaned = cpp.transform(xtrain)
      file_X_train_cleaned_sparse_embed = data_folder / \
          'x_train_cleaned_sparse_embed_full.pkl'
      joblib.dump(X_train_cleaned, file_X_train_cleaned_sparse_embed)
```

```
/content/drive/MyDrive/NLP/custom-functions/custom_preprocessor_mod.py:90:
MarkupResemblesLocatorWarning: The input looks more like a filename than markup.
You may want to open this file and pass the filehandle into BeautifulSoup.
    soup = BeautifulSoup(text, "html.parser")
```

```
[19]: ['/content/drive/MyDrive/NLP/datasets/spam/x_train_cleaned_sparse_embed_full.pkl']
```

```
[20]: X_test_cleaned = cpp.transform(xtest)
      # save this to a file
      file_X_test_cleaned_sparse_embed = data_folder /
      ↪ 'x_test_cleaned_sparse_embed_full.pkl'
      joblib.dump(X_test_cleaned, file_X_test_cleaned_sparse_embed)
```

```
[20]: ['/content/drive/MyDrive/NLP/datasets/spam/x_test_cleaned_sparse_embed_full.pkl']
```

0.4 Defining Class Weights for Imbalanced data to train Classifier

```
[21]: w = {}

      w[1] = int(y_train.value_counts()[0]/ytrain.shape[0]*100)

      w[0] = 100 - w[1]

      print(w)
```

```
{1: 86, 0: 14}
```

0.5 Final Pipeline (Pipeline 3): Combine Manual Features and TfID vectors

```
[ ]:
```

```
[22]: featurizer = ManualFeatures(spacy_model='en_core_web_sm')
```

```
[23]: X_train_features, feature_names = featurizer.fit_transform(xtrain)
```

```
/content/drive/MyDrive/NLP/custom-functions/custom_preprocessor_mod.py:90:
MarkupResemblesLocatorWarning: The input looks more like a filename than markup.
You may want to open this file and pass the filehandle into BeautifulSoup.
    soup = BeautifulSoup(text, "html.parser")
```

```
[24]: print(X_train_features.shape)
      X_train_features[0:3]
```

```
(4457, 11)
```

```
[24]: array([[13.      , 58.      , 46.      , 3.28571429, 0.      ,
           0.      , 0.      , 4.      , 1.      , 3.      ,
           0.      ],
          [13.      , 51.      , 39.      , 2.78571429, 0.      ,
           0.      , 0.      , 1.      , 3.      , 3.      ,
           0.      ],
          [10.      , 53.      , 44.      , 4.      , 0.      ,
           0.      , 0.      , 3.      , 0.      , 1.      ,
           1.      ]])
```

```
[25]: X_train_cleaned_sparse_embed = joblib.load(file_X_train_cleaned_sparse_embed)

X_train_final = pd.concat((pd.DataFrame(X_train_cleaned_sparse_embed,
    ↪columns=['cleaned_text']),
                           pd.DataFrame(X_train_features,
    ↪columns=feature_names)), axis=1)

X_train_final.head()
```

```
[25]:
```

	cleaned_text	count_words	count_characters	\
0	sen tell go join uncle finance cbe	13.0	58.0	
1	reach home n bathe liao u	13.0	51.0	
2	hello site download song urgent pls	10.0	53.0	
3	hi:)ct employee	5.0	28.0	
4	aight pick open tonight	10.0	49.0	

	count_characters_no_space	avg_word_length	count_digits	count_numbers	\
0	46.0	3.285714	0.0	0.0	
1	39.0	2.785714	0.0	0.0	
2	44.0	4.000000	0.0	0.0	
3	24.0	4.000000	0.0	0.0	
4	40.0	3.636364	0.0	0.0	

	noun_count	aux_count	verb_count	adj_count	ner
0	0.0	4.0	1.0	3.0	0.0
1	0.0	1.0	3.0	3.0	0.0
2	0.0	3.0	0.0	1.0	1.0
3	0.0	1.0	1.0	1.0	0.0
4	1.0	2.0	1.0	2.0	0.0

```
[26]: class SparseTransformer(TransformerMixin, BaseEstimator):
        def __init__(self):
            pass

        def fit(self, X, y=None):
            return self
```

```

def transform(self, X, y=None):
    return csr_matrix(X)

sparse_features = Pipeline([('sparse', SparseTransformer()), ])
vectorizer = Pipeline([('tfidf', TfidfVectorizer(max_features=5)), ])

combined_features = ColumnTransformer(
    transformers=[
        ('tfidf', vectorizer, 'cleaned_text'),
    ], remainder=sparse_features
)

```

0.5.1 Create Final Pipeline

```

[27]: classifier_3 = Pipeline([('combined_features', combined_features),
                              ('classifier', LogisticRegression(max_iter=10000,
                              ↪random_state=21, class_weight=w)),
                              ])

```

```

[53]: # We are exploring a small combination of parameters
      # If the search space is very large then we should use RandomSerachCV or some
      ↪other methods

param_bayes_classifier_3 = {'combined_features__tfidf__tfidf__max_features':
      ↪Integer(500, 10000),
                           #'combined_features__tfidf__tfidf__ngram_range':
      ↪Categorical([(1, 1), (1, 2), (1, 3)]),
                           'combined_features__tfidf__tfidf__max_df': Real(0.2,
      ↪0.8),
                           'combined_features__tfidf__tfidf__min_df': Real(0.
      ↪01, 0.05, prior='log-uniform'),
                           'classifier__solver': Categorical(['liblinear',
      ↪'saga', 'newton-cg', 'lbfgs']),
                           'classifier__C': Real(0.001, 1000,
      ↪prior='log-uniform')
                           }

```

```

[54]: # Define a custom scoring function for PR AUC
def custom_pr_auc_scorer(y, y_proba):
    #y_proba = estimator.predict_proba(X)[: , 1] # Probability of positive class
    precision, recall, _ = precision_recall_curve(y, y_proba)
    pr_auc = auc(recall, precision)
    return pr_auc

```



```

# Define a custom scoring function for Cohen's Kappa
def custom_kappa_scorer(y, y_pred):
    kappa = cohen_kappa_score(y, y_pred)
    return kappa

# Use cross-validation with the custom scoring function
pr_auc_scorer = make_scorer(custom_pr_auc_scorer, greater_is_better=True) #
↳ Set greater_is_better=True for higher PR AUC scores

# Use cross-validation with the custom scoring function
kappa_scorer = make_scorer(custom_kappa_scorer)

```

0.5.2 Perform Bayesian Optimization

```

[55]: # Perform Bayesian optimization
optimizer3 = BayesSearchCV(estimator=classifier_3, search_spaces=
↳ param_bayes_classifier_3, n_iter=66, cv=5, n_jobs = -1, verbose = 1
, scoring=pr_auc_scorer
)

# Fit the model to the data
optimizer3.fit(X_train_final, ytrain)

# Print the best hyperparameters aand corresponding score
print("Best score:", optimizer3.best_score_)
print("Best hyperparameters:", optimizer3.best_params_)

```

```

Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits

```



```
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Fitting 5 folds for each of 1 candidates, totalling 5 fits
Best score: 0.931513850007699
Best hyperparameters: OrderedDict([('classifier__C', 0.17159080752918784),
('classifier__solver', 'newton-cg'), ('combined_features__tfidf__tfidf__max_df',
0.8), ('combined_features__tfidf__tfidf__max_features', 10000),
('combined_features__tfidf__tfidf__min_df', 0.01)])
```

```
[56]: print("Best cross-validation score: {:.2f}".format(optimizer3.best_score_))
print("\nBest parameters: ", optimizer3.best_params_)
print("\nBest estimator: ", optimizer3.best_estimator_)
```

```
Best cross-validation score: 0.93
```

```
Best parameters:  OrderedDict([('classifier__C', 0.17159080752918784),
('classifier__solver', 'newton-cg'), ('combined_features__tfidf__tfidf__max_df',
0.8), ('combined_features__tfidf__tfidf__max_features', 10000),
('combined_features__tfidf__tfidf__min_df', 0.01)])
```

```
Best estimator:  Pipeline(steps=[('combined_features',
ColumnTransformer(remainder=Pipeline(steps=[('sparse',
SparseTransformer())])),
transformers=[('tfidf',
Pipeline(steps=[('tfidf',
TfidfVectorizer(max_df=0.8,
max_features=10000,
min_df=0.01))])),
('cleaned_text')]))),
('classifier',
LogisticRegression(C=0.17159080752918784,
class_weight={0: 14, 1: 86}, max_iter=10000,
random_state=21, solver='newton-cg'))])
```

0.5.3 Save & Load Model

```
[57]: file_best_estimator_pipeline3 = model_folder / \
'logistic_final.pkl'
file_complete_bayes_pipeline3= model_folder / \
'logistic_final_complete_bayes_full.pkl'

joblib.dump(optimizer3.best_estimator_, file_best_estimator_pipeline3)
joblib.dump(optimizer3, file_complete_bayes_pipeline3)


# load the saved model
best_estimator_pipeline3_round1 = joblib.load(
file_best_estimator_pipeline3)
complete_bayes_pipeline3_round1 = joblib.load(
```

```
file_complete_bayes_pipeline3)
```

0.5.4 Plot Learning Curve

```
[58]: # plot learning curves
plot_learning_curve(best_estimator_pipeline3_round1, 'Learning Curves',
                    X_train_final, ytrain, n_jobs=-1)
```

```
[58]: <module 'matplotlib.pyplot' from '/usr/local/lib/python3.10/dist-
packages/matplotlib/pyplot.py'>
```



```
[59]: # let's check the train scores
print(best_estimator_pipeline3_round1.score(X_train_final, y_train))

# let's check the cross validation score
print(complete_bayes_pipeline3_round1.best_score_)
```

```
0.9847431007404084
```

```
0.931513850007699
```

```
[60]: X_train_final.shape
```

```
[60]: (4457, 12)
```

```
[60]:
```

0.5.5 Evaluate on Test

```
[61]: # Final Pipeline
def final_pipeline(text):
    cleaned_text = cpp.transform(text)
    # cleaned_text = joblib.load(file_X_test_cleaned_sparse_embed)
    X_features, feature_names = featurizer.fit_transform(text)
    X_final = pd.concat((pd.DataFrame(cleaned_text, columns=['cleaned_text']),
```

```

pd.DataFrame(X_features, columns=feature_names)),  

axis=1)  

best_estimator_pipeline3_round1 = joblib.load(  

    file_best_estimator_pipeline3)  

predictions = best_estimator_pipeline3_round1.predict(X_final)  

return predictions

```

```

[62]: # predicted values for Test data set  

y_test_pred = final_pipeline(xtest)

```

```

/content/drive/MyDrive/NLP/custom-functions/custom_preprocessor_mod.py:90:
MarkupResemblesLocatorWarning: The input looks more like a filename than markup.
You may want to open this file and pass the filehandle into BeautifulSoup.
soup = BeautifulSoup(text, "html.parser")

```

```
[62]:
```

```

[63]: print('\nTest set classification report:\n\n',  

classification_report(y_test, y_test_pred))

```

Test set classification report:

	precision	recall	f1-score	support
0	0.99	0.98	0.98	966
1	0.87	0.93	0.90	149
accuracy			0.97	1115
macro avg	0.93	0.96	0.94	1115
weighted avg	0.97	0.97	0.97	1115

```

[64]: # prompt: plot confusion matrix  

print(confusion_matrix(ytest, y_test_pred))

```

```

[[945  21]
 [ 10 139]]

```

```
[65]: print(custom_pr_auc_scorer(ytest, y_test_pred))
```

```
0.9053022579528697
```

```
[66]: print(balanced_accuracy_score(ytest, y_test_pred))
```

```
0.9555733878027429
```

0.5.6 Final Score on the Chosen Metric

Precision Recall AUC - 0.9019822478991

[]: